

Neural Implicit Flow: A Comparative Study

Implementation Approaches and Performance Analysis

Christian Beneke

10.02.2025

Outline

- 1 Introduction
- 2 Neural Implicit Flow Framework
- 3 Implementation Approaches
- 4 Experimental Setup
- 5 Results and Analysis
- 6 Conclusions

Problem Space & Motivation

- High-dimensional spatio-temporal dynamics are computationally challenging
- Traditional methods face limitations:
 - Fixed mesh structures (SVD, CAE)
 - Poor scaling with high-dimensional data
 - Limited ability to capture nonlinear dynamics
- Real-world applications require:
 - Adaptive mesh handling
 - Efficient dimensionality reduction
 - Real-time processing capabilities

Key Challenges in Current Approaches

- Singular Value Decomposition (SVD)
 - Limited to fixed mesh structures
 - Poor scaling with dimensionality
- Convolutional Autoencoders (CAE)
 - Requires uniform grid resolution
 - Struggles with adaptive meshing
- Common Issues
 - High computational overhead
 - Limited expressiveness for complex dynamics
 - Poor generalization to new scenarios

- Two specialized networks:
 - ShapeNet: Spatial complexity
 - ParameterNet: Temporal evolution
- Key innovations:
 - Mesh-agnostic representation
 - Efficient parameter sharing
 - Adaptive complexity handling




Figure: NIF Architecture

- Core mapping function:

$$f_{\theta} : (\mathbf{x}, \mathbf{t}) \mapsto u(\mathbf{x}, \mathbf{t})$$

where:

- $\mathbf{x} \in \mathbb{R}^d$: Spatial coordinate
 - $\mathbf{t} \in \mathbb{R}^p$: Temporal/parametric input
 - u : Field value
- Hypernetwork decomposition:

$$f_{\theta}(\mathbf{x}, \mathbf{t}) = \text{ShapeNet}_{\text{ParameterNet}_{\theta_p}(\mathbf{t})}(\mathbf{x})$$

Implementation Overview

- Three distinct implementations:
 - Upstream Reference (TensorFlow)
 - TensorFlow Functional API
 - PyTorch Implementation
- Key considerations:
 - Framework-specific optimizations
 - Code maintainability
 - Development ergonomics

Upstream Reference Implementation

- Based on original paper implementation
- Key modifications:
 - Migration to modern TensorFlow patterns
 - Improved gradient tape management
 - Enhanced base class structure

TensorFlow Functional API Implementation

- Complete architectural redesign
- Key features:
 - Functional programming principles
 - Improved composability
 - XLA compilation support
- Optimizations:
 - Vectorized weight generation
 - Memory-efficient parameter handling
 - Mixed precision training (FP16/FP32)

- Modern PyTorch features:
 - Dynamic computation graphs
 - Native autograd functionality
 - Improved debugging capabilities
- Implementation highlights:
 - Flexible activation mapping
 - Efficient static weight layers
 - Comprehensive training logger

Low Frequency Wave

- Domain: $x \in [0, 1]$, $t \in [0, 100]$
- 200 spatial points, 10 timesteps
- Simple periodic traveling wave
- $u(x, t) = \exp(-1000(x - x_0 - ct)^2)$




Figure: Low Frequency Example

High Frequency Wave

- Same spatial/temporal domain
- Added frequency $\omega = 400$
- Complex oscillatory pattern
- $u(x, t) = \exp(-1000(x - x_0 - ct)^2) \sin(\omega(x - x_0 - ct))$




Figure: High Frequency Example

ParameterNet

- Dense input layer
- Two hidden layers (30 units)
- Single-unit bottleneck
- Adaptive output layer

ShapeNet Variants

- Shortcut (Low frequency)
 - Skip connections
 - Swish activation
- SIREN (High frequency)
 - Sinusoidal activation
 - $\omega_0 = 30$ scaling

- Low Frequency Results:
 - TF Functional API (Adam): $1.526e-04$
 - PyTorch (Adam): $1.549e-04$
 - Upstream (AdaBelief): $5.073e-04$
- High Frequency Results:
 - PyTorch (Adam): $1.884e-04$
 - TF Functional API (Adam): $5.415e-04$
 - Upstream: Not available

Processing Speed Comparison

Implementation	Low Freq.	High Freq.
TF Functional	11.5-17K pts/s	11.5-16K pts/s
PyTorch	12-15K pts/s	6-16K pts/s
Upstream	12-17K pts/s	-

Table: Processing Speed (points per second)

- Adam vs AdaBelief comparison:
 - TF Functional API: 62.5% improvement
 - PyTorch: 5.10x improvement (low freq.)
 - PyTorch: 11.77x improvement (high freq.)
- Key findings:
 - Adam: More stable convergence
 - Better final loss values
 - Consistent across implementations

Visual Results - Low Frequency




Figure: Low Frequency Predictions (TF Functional API)

Visual Results - High Frequency




Figure: High Frequency Predictions (TF Functional API)

- Implementation Trade-offs:
 - Modern implementations excel at high-frequency cases
 - TF Functional API: Most consistent performance
 - PyTorch: Best high-frequency accuracy
- Practical Implications:
 - All implementations achieve production-ready speed
 - Adam optimizer consistently outperforms AdaBelief
 - Framework choice impacts development experience

- Technical Improvements:
 - Additional network architectures
 - Memory optimization
 - Computational efficiency
- Application Areas:
 - Fluid dynamics
 - Climate modeling
 - Materials science
- Research Opportunities:
 - Complex spatio-temporal problems
 - Multi-scale phenomena
 - Real-time applications

Thank you for your attention!

Questions?