

# Neural Implicit Flow: A Comparative Study

## Implementation Approaches and Performance Analysis

Christian Beneke

10.02.2025

# Outline

- 1 Introduction
- 2 Neural Implicit Flow Framework
- 3 Implementation Approaches
- 4 Experimental Setup
- 5 Results and Analysis
- 6 Conclusions

# Problem Space & Motivation

- High-dimensional spatio-temporal dynamics are computationally challenging
- Traditional methods face limitations:
  - Fixed mesh structures (SVD, CAE)
  - Poor scaling with high-dimensional data
  - Limited ability to capture nonlinear dynamics
- Real-world applications require:
  - Adaptive mesh handling
  - Efficient dimensionality reduction
  - Real-time processing capabilities

# Core Architecture

Two specialized networks:

- ShapeNet: Spatial complexity
- ParameterNet: Temporal evolution

Key innovations:

- Mesh-agnostic representation
- Efficient parameter sharing
- Adaptive complexity handling




Figure: NIF Architecture

- Core mapping function:

$$f_{\theta} : (\mathbf{x}, \mathbf{t}) \mapsto u(\mathbf{x}, \mathbf{t})$$

where:

- $\mathbf{x} \in \mathbb{R}^d, d \in \mathbb{N}$ : Spatial coordinate
- $\mathbf{t} \in \mathbb{R}^p, p \in \mathbb{N}$ : Temporal/parametric input
- Hypernetwork decomposition:

$$f_{\theta}(\mathbf{x}, \mathbf{t}) = \text{ShapeNet}_{\text{ParameterNet}_{\theta_p}(\mathbf{t})}(\mathbf{x})$$

# Implementation Overview

- Three distinct implementations:
  - Upstream Reference (TensorFlow)
  - TensorFlow Functional API
  - PyTorch Implementation
- Key considerations:
  - Framework-specific optimizations
  - Code maintainability

# Common Setup

- Implemented base class for common functionality
- Implemented two of the given example cases
- Made optimiser configurable (Adam, AdaBelief)
- Comprehensive training logger with visualization

# Upstream Reference Implementation

- Based on original paper implementation
- Key modifications:
  - Migration to current TensorFlow version
  - Various bug fixes needed to run the provided examples
- Implementation details:
  - Direct TensorFlow.Keras Model inheritance
  - Static layer definitions with fixed weights
  - Manual parameter mapping and weight updates
  - Internal forward pass in a single function call



# Upstream Reference Architecture

Layer (type)	Output Shape	Param #
first_dense_pnet (Dense)	(None, 30)	60
hidden_mlpshortcut_pnet_0 (MLP_SimpleShortCut)	(None, 30)	930
hidden_mlpshortcut_pnet_1 (MLP_SimpleShortCut)	(None, 30)	930
bottleneck_pnet (Dense)	(None, 1)	31
last_pnet (Dense)	(None, 1951)	3,902

Total params: 5,853 (22.86 KB)

Figure: Upstream Reference Model Summary

# TensorFlow Functional API Implementation

- Complete architectural redesign
- Key features:
  - Utilises TensorFlow's functional API
  - Custom SIREN kernel initialization
  - Dynamic layer generation
- Optimizations:
  - Automatic shape inference and building
  - TF Function decoration for performance
  - Full ParameterNet output forwarded to all ShapeNet layers

# TensorFlow Functional API Architecture

Layer (type)	Output Shape	Param #
first_pnet (Dense)	(None, 30)	60
hidden_pnet_0 (Shortcut)	(None, 30)	930
hidden_pnet_1 (Shortcut)	(None, 30)	930
bottleneck_pnet (Dense)	(None, 1)	31
last_pnet (Dense)	(None, 1951)	3,902
first_snet (HyperDense)	(None, 30)	0
hidden_snet_0 (HyperDense)	(None, 30)	0
hidden_snet_1 (HyperDense)	(None, 30)	0
last_snet (HyperDense)	(None, 1)	0

Total params: 5,853 (22.86 KB)

Figure: TensorFlow Functional API Model Summary

- Core Features:
  - Mixed precision training support (FP16/BF16)
  - Custom StaticDense and StaticSIREN layers
  - Gradient scaling and automatic mixed precision (AMP)
- Implementation Details:
  - Custom AdaBelief optimizer with gradient centralization
  - Configurable learning rate scheduler with warmup
  - Dynamic device handling (CPU/GPU) with dtype management

## Low Frequency Traveling Wave

- Domain:  $x \in [0, 1]$ ,  $t \in [0, 100]$
- 200 spatial points, 10 timesteps
- Simple periodic traveling wave
- $u(x, t) = \exp(-1000(x - x_0 - ct)^2)$




Figure: Low Frequency Example

## High Frequency Traveling Wave

- Same spatial/temporal domain
- Added frequency  $\omega = 400$
- Complex oscillatory pattern
- $u(x, t) = \exp(-1000(x - x_0 - ct)^2) \sin(\omega(x - x_0 - ct))$




Figure: High Frequency Example

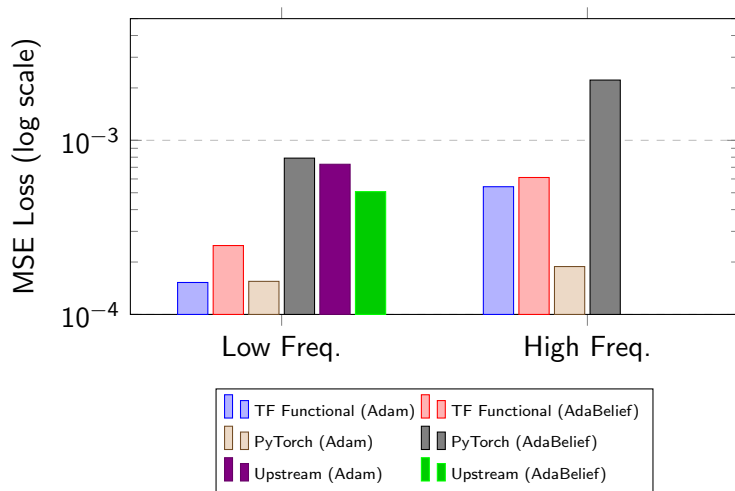
## ParameterNet

- Dense input layer
- Two hidden layers (30 units)
- Bottleneck layer (1 unit)
- Adaptive output layer

## ShapeNet Variants

- Shortcut (Low frequency)
  - Skip connections
  - Swish activation
- SIREN (High frequency)
  - Sinusoidal activation
  - $\omega_0 = 30$  scaling

# Performance Overview





- Adam vs AdaBelief comparison:
  - TF Functional API: 62.5% improvement
  - PyTorch: 5.10x improvement (low freq.)
  - PyTorch: 11.77x improvement (high freq.)
- Key findings:
  - Adam: More stable convergence
  - Better final loss values
  - Consistent across implementations

# Visual Results - Low Frequency




Figure: Low Frequency Predictions (TF Functional API)

# Visual Results - High Frequency




Figure: High Frequency Predictions (TF Functional API)

# Key Findings

- Implementation Trade-offs:
  - Modern implementations have consistent performance
  - PyTorch: Best high-frequency accuracy
- Practical Implications:
  - All implementations achieve production-ready speed
  - Adam optimizer consistently outperforms AdaBelief

# Planned vs Achieved Goals

- ✓ Port from TensorFlow to PyTorch
- + Add TensorFlow Functional API implementation
- ✓ Implement using upstream SIREN
- ✓ Improve parallelism in training phase
- × Compare to other HyperNetworks / LoRA / etc
- ✓ Compare existing results with own implementation
- × Compare performance with very hard example (e.g. weather prediction)

# Thank you for your attention!

Questions?